Tesina

Angelica Urbanelli S271114

Data Spaces, a.y. 2020/2021

Contents

1   Introduction

-spiegare il problema importanza di distinguere buoni e cattivi pagatori statistiche?
   -interpretabilità? -falsi buoni peggio di falsi cattivi
   -tutto il codice si trova nella repo github xx -per navigarlo e browsing all images usare il link xx di jupyternotebook

2   Dataset description

2.1   History

The dataset used comes from the UCI Machine Learning Repository [5], under the name "South German Credit (UPDATE) Data Set" [8].

Ulrike Grömping, professor at the Beuth University in Berlin, in her paper [7] provides the history of this dataset, her considerations about the data and corrections on the code table.

Basically, the data come from a large regional bank in the southern Germany that have been collected from 1973 to 1975, and have been originally provided to UCI in 1994 by Professor Dr. Hans Hofmann from Hamburg University [2] as part of a group of datasets in the context of the EU Statelog Project.

Because of many inconsistencies, found while trying to interpret the final results of her experiments, Grömping decided to research the story of this data, that she found in the German literature together with the same dataset with some differences. These informations helped her to fix the code table (a file that explains the encoding of categorical variables) of this dataset and consequently to provide the correct one (now attached in the .zip downloadable from UCI).

Grömping also explained that it was worth it because, although the dataset contains very old data, it is widely used in many researches in the domain of interpretable machine learning, indeed there are various R packages that include this data. In addition, it *is one of the few data sets on credit scoring that has a meaning attached to variables and their levels*, which is a very important feature when using this kind of data to do experiments whose interpretability is a key point of research.

2.2   General structure

The dataset contains 1000 samples, each one characterized by 20 features and classified as `good` or `bad credit risk`, in particular there are 700 good ones and 300 bad ones [Figure 1]. Customers with good credits perfectly complied with the conditions of the contract, while customers with bad credits did not comply with the contract as required. As reported in the aforementioned paper [7], the actual percentage of bad credits was around 5%, and examples of bad credit risk have been heavily oversampled.


Figure 1: `good` and `bad` credit risk distribution

2.3   Features description

Among the 20 features, there are 3 numerical discrete variables:

- **duration**: credit duration in months

- **amount**: credit amount in DM[1]; original values are not available, the ones present in the dataset are the result of an unknown monotonic transformation

- **age**: age of the debtor in years

10 ordinal variables; most of theme were numerical ones on which binning has been applied, that means that they have been aggregated into a fixed number of intervals, so that they can be treated as ordinal features:

- **employment_duration**: duration of debtor's employment with current employer (unemployed; < 1 year ; ≥ 1 and < 4 years; ≥ 4 and < 7 years; ≥ 7 years)

- **installment_rate**: credit installments expressed as a percentage of debtor's disposable income (≥ 35; ≥ 25 and < 35; ≥ 20 and < 25; < 20); it is the only ordinal feature expressed in a decreasing order

- **present_residence**: from how many years the debtor lives in the present residence (< 1 year ; ≥ 1 and < 4 years; ≥ 4 and < 7 years; ≥ 7 years)

- **number_credits**: number of credits including the current one the debtor has (or had) at this bank (1; 2 or 3; 4 or 5; ≥ 6)

- **people_liable**: number of people who financially depend on the debtor (i.e., are entitled to maintenance) (from 0 to 2; 3 or more)

- **status**[2]: status of the debtor's checking account with the bank in DM (no checking account; < 0; $0 \leq ... < 200$; ≥ 200 / salary for at least 1 year)

- **savings**[2]: debtor's savings in DM (unknown/no savings account; < 100; $100 \leq ... < 500$; $500 \leq ... < 1000$; ≤ 1000)

- **credit_history**[2]: history of compliance with previous or concurrent credit contracts (delay in paying off in the past; critical account/other credits elsewhere; no credits taken/all credits paid back duly; existing credits paid back duly till now; all credits at this bank paid back duly)

- **job**: quality of debtor's job (unemployed/unskilled - non-resident; unskilled - resident; skilled employee/official; manager/self-employed/highly qualified employee)

- **property**: the debtor's most valuable property, i.e. the highest possible code is used (unknown / no property; car or other [savings don't fall into this category]; building society savings agreement (mortgage)/life insurance; real estate)

7 categorical variables:

- **purpose**: purpose for which the credit is needed (others; car (new); car (used); furniture/equipment; radio/television; domestic appliances; repairs; education; vacation; retraining; business)

- **personal_status_sex**: combined information on sex and marital status; sex cannot be recovered from the variable because male singles and female non-singles are coded with the same code; in addition, female widows are not listed in any of the categories (male divorced/separated; female non-single or male single; male married/widowed; female single)

- **other_debtors**: whether there is another debtor or a guarantor for the credit (none; co-applicant; guarantor)

- **other_installment_plans**: installment plans from providers other than the credit-giving bank (bank; stores; none)

- **housing**: type of housing the debtor lives in (for free; rent; own)

- **telephone**: whether there is a telephone landline registered on the debtor's name; of course this variable would have no meaning nowadays, but this data come from 1970s (yes; no)

- **foreign_worker**: whether the debtor is a foreign worker (yes; no)

---

[1]stands for Deutsche Mark, was the official currency of West Germany from 1948 until 1990 and later the unified Germany from 1990 until 2002 [14]

[2]Those are considered as categorical by Grömping [7], but in my opinion their labels can be ranked and also considering that to properly manage categorical features an encoding have to be done, thus likely this brings to a huge number of features. Possibly, the position of *no checking account* in status feature could be discussed with a domain expert.

3   Dataset Analysis

Figure 2: Features' distributions

3.1   Features' distributions

Here [Figure 2] the distribution of the various features can be seen. Some of them are highly imbalanced, for instance it can be noticed that almost all customers are not foreigners and have neither another debtor nor a guarantor for the credit.

It might be interesting to take a look at the box plots of the numerical discrete features: `duration`, `amount` ad `age` [Figure 3]. Given that Q_1 and Q_3 are, respectively, the first and the third quartile, and that the interquartile range $IQR = Q_3 - Q_1$, in the chosen representation the whiskers are: the largest observed point that falls within Q_3 and $Q_3 + 1.5 \cdot Q_3$ and the lowest observed point that falls within Q_1 and $Q_1 - 1.5 \cdot Q_1$ [13]. The single points, instead, are highlighted in red if they fall within the lowest whisker and $4 \cdot Q_1 - 3 \cdot Q_3$ or within the highest whisker and $4 \cdot Q_3 - 3 \cdot Q_1$ [12], those are called *suspected* outliers; while the blue points are the ones outside these ranges, thus they can be considered outliers beyond any doubt.

In this case, since those plots represent the univariate distribution, those points are not considered as outliers; they will be better evaluated and managed in section 4.3.

Figure 3: Box Plots of `duration`, `amount` ad `age`

3.2   Features distribution per class

In figure [4] and [5], the same plots of section 3.1 are shown, highlighting the distributions separately for the two class labels.

4   Data preprocessing

4.1   Labels encoding

All the features' labels are expressed as integers: both categorical and ordinal ones have been previously mapped to integers by the donor of the data, using label encoding either from 0 to N-1 (only for **purpose** and **credit__history**) or from 1 to N (all the other variables) where N is the number of labels for a certain feature. Apparently there is no particular reason for this distinction. For what concerns class labels (`credit risk`), they have been mapped to 0 for `bad` and 1 for `good`.

In order to make variables more uniform, two small changes have been done:

- the variable **installment__rate** is the only one among the ordinal variables having a decreasing order, so its mapping has been inverted; thus now label 1 means < 20, label 2 means ≥ 20 and < 25 and so on;

- **credit__history** and **purpose**'s mappings have been changed from $0 \to N-1$ to $1 \to N$;

- in this kind of classification, the *positive* class is the one we are more interested in correctly classifying, in this case the `bad` creditors. By convention, the *positive* class has label 1, and the other, the *negative* one, has label 0; thus, the class labels have been inverted.

In addition, all categorical variables have to be encoded in a different way with respect to the actual one $(1 \to N)$. Indeed this encoding gives an arbitrary ranking to features that do not have one, and this is a problem when applying distance-based classification algorithms because those could potentially exploit this fictional structure created by the mapping itself.

4.2   One hot encoding vs binary encoding

In order to perform this mapping, two encoders have been taken into account

One hot encoding   For each unique value in a variable, a new column is created, whose values are either 1s or 0s, depending on whether the value matches the column header.
It is very simple, it does not interfere with interpretability since every new column has a specific meaning, and it allows very well to separate categorical features' labels. However, the downside is that we may end up with a huge number of features, especially if we need to map variables with an high number of labels.

Figure 4: Features' distributions per class

Binary encoding Values are firstly converted into their binary code, and then the digits are split into separate columns.

With this method, the overhead due to the increase in the number of features is limited, because a feature with N distinct values is mapped into $\log_2 N$ columns instead of N. The drawbacks are that the resulting features still have a weak binding between them and the new columns do not really have a specific meaning on their own.

Given these considerations, among the categorical variables:

- **purpose** (11 labels) has been encoded with binary encoding;

- **other_debtors**, **other_installment_plans**, **housing** and **personal_status_sex** (respectively with 3, 3, 3 and 4 labels) have been encoded using one hot encoding, since the benefit of binary encoding is not worthwhile with these small number of labels;

- **telephone** and **foreign_worker** have only 2 labels, therefore do not need any encoding.

Moreover, since one of the 3 labels of **other_debtors** is `none`, it can be encoded with only two columns, where `none` is encoded with both at 0.
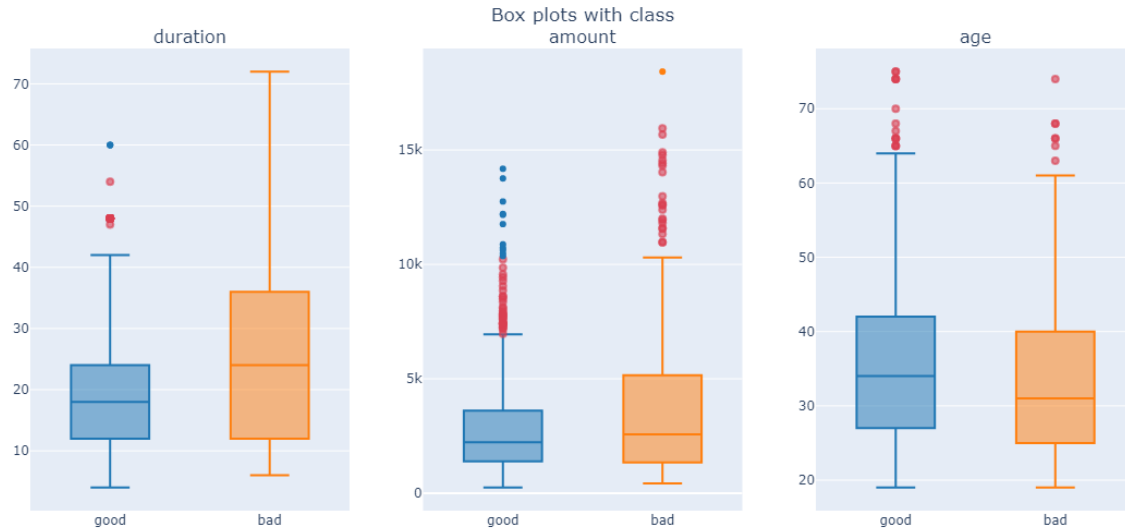
Figure 5: Box Plots of `duration`, `amount` ad `age` divided by class

4.3 Missing values and outliers

The dataset does not contain any missing value.

Regarding outliers detection, by looking at figures [3] and [5] (for more boxplots see:) it might seem that the dataset has a huge number of outliers; however, observing one feature at a time can be misleading, since the whole situation should be taken into account. For this reason, instead of considering the univariate distributions separately, can be a good idea to perform a multivariate outliers detection.

For this purpose, various distance metrics and techniques exist. Among them, Mahalanobis Distance (introduced by Mahalanobis in 1936 [1]) is an effective one, its goal is to find the distance between a point $\vec{x}$ and a distribution. Differently from other techniques, it can manage distributions where every feature has a different scale and variance. Furthermore, by using the covariance matrix, it is able to detect outliers basing on the distribution of points, unlike e.g. the Euclidean distance [Figure 6].
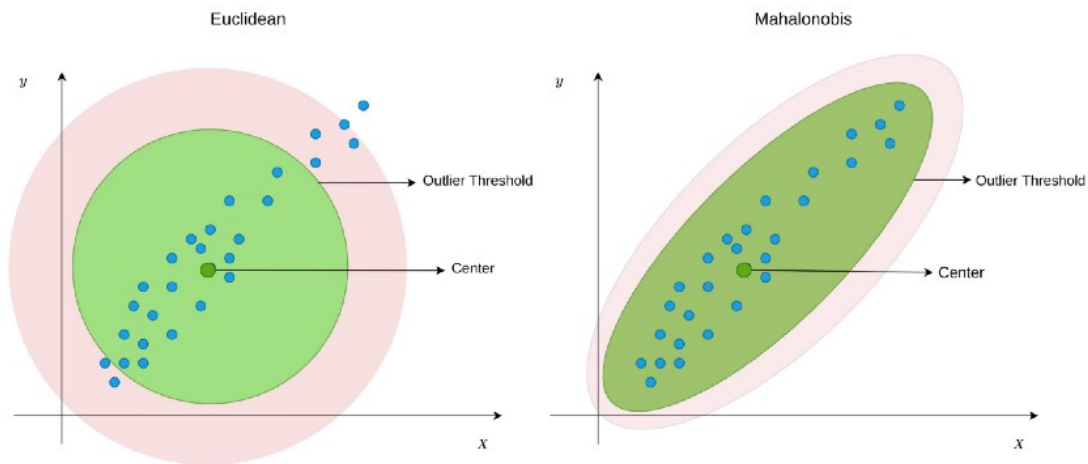

Figure 6: Euclidean distance vs Mahalonobis distance, image by [10]

Given a point $\vec{x} = (x_1, x_2, ..., x_N)^T$ extracted from a distribution of m points in $\vec{X} = (X_1, X_2, ..., X_N)^T$ where

Figure 7: Box plot of Mahalonobis distance distribution

Figure 8: Violin plot of Mahalonobis distance distribution

$X_1, X_2, ..., X_N$ are the random variables of the dataset; given $\vec{\mu} = (\mu_1, \mu_2, ..., \mu_N)^T$ the mean of the set of observations whose entries $\mu_i = \frac{1}{m}\sum x_i$, and the covariance matrix Σ whose entry

$$\Sigma_{(i,j)} = cov[X_i, X_j] = E[(X_i - E[X_i])(X_j - E[X_j])] \tag{1}$$

thus, the Mahalanobis Distance from a point $\vec{x}$ and the set of points it has been extracted from is

$$d_M(\vec{x}, \vec{\mu}) = \sqrt{(\vec{x} - \vec{\mu})^T \Sigma^{-1} (\vec{x} - \vec{\mu})} \tag{2}$$

The covariance matrix (equation 1) is estimated using the `numpy.cov` function, see the documentation [11] and the code for better understanding.

Once computed the Mahalanobis distance for every point, its distribution can be evaluated to detect multidimensional outliers. As can be seen in figures [7] and [8], there are a few points (seven) whose distance is greater than $Q_3 + 1.5 \cdot Q_3$, thus they could be considered as multidimensional outliers. However, they are a very small number of points and, among those, four have class label 0 (bad credit risk), that is the one with the lowest number of observations. That means that they could seem outliers because there are not so many samples of this label (considering that the ones present have already been oversampled), thus they could be significant points in detecting bad credit risk. In addition, all those points fall in the range of the so called *suspected outliers*, while there are no points $< 4 \cdot Q_1 - 3 \cdot Q_3$ or $> 4 \cdot Q_3 - 3 \cdot Q_1$ (that are the ranges where points are very likely to be outliers). For these reasons, those seven points are not removed from the dataset.

4.4   Data normalization

When performing analysis with datasets whose features have different scales, data normalization is a very important step in the preprocessing procedure. Indeed, especially when applying some specific techniques that rely on distance or variance (e.g. distance-based classification algorithm, PCA and so on), if features are not rescaled, the distance between points and data variance are more affected by those features that have larger scale or higher values. Thus, the features end up not to have the same importance.

In addition, although some algorithms' implementations take care of centering data (e.g. `sklearn.decomposition` `.PCA`[3]), others do not (e.g. `sklearn.decomposition.KernelPCA`), thus in order to avoid problems in data analysis, data points are also centred.

Specifically, the following transformation has been applied to features, that is, for each feature X, the new one $\hat{X}$:

$$\hat{X} = \frac{X - mean(X)}{max(X) - min(X)} \tag{3}$$

---

[3]documentation at: `https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html`
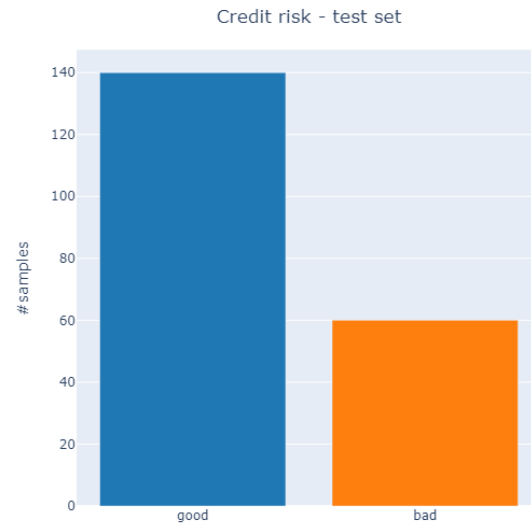
Figure 9: Training set

Figure 10: Test set

in this way, every feature is centered ($\mu_i = 0$) and rescaled in a unitary range.

4.5   Training-test split

To proceed in the classification phase, the dataset has been split into training and test set respectively for 80% and 20% in a stratified way, that means that in both sets, the proportion of `good` and `bad` credit risk of the original dataset has been kept [Figure 9 and 10]. This is particularly important, because the test set (the one on which the trained classification algorithm will be evaluated), should be as much similar as possible to real data, that means keeping the proportion of label distribution, do not applying any further transformation to those data and do not use them during training. Indeed, all the dimensionality reduction techniques presented in the next section (5) has been applied only using the training set.

5   Dimensionality reduction

5.1   Curse of dimensionality

After the encoding phase, the number of features (label excluded) went from 20 to 31. Of course, the encoding was needed, but in data analysis having to deal with many features can be problematic for a lot of reasons. Indeed this is called *the curse of dimensionality*. When the dimensions increase, the volume of the hypercube containing the data increases as well, so all the points become more distant between them, thus making the dataset very sparse. To overcome this problem, the number of data points should increase exponentially with the dimension. Basically, all points become distant and at the same time they have similar distances among them; therefore, the concepts of nearest neighbours and distance become pointless. So, in cases like this, it is always a good idea to consider using a dimensionality reduction technique.

5.2   Correlation based

One first intuitive way to reduce dimensionality is by looking at the pairwise correlation between features, in order to spot possible duplicated ones. In this case [Figure 11], most of the features are almost not correlated, or with very low values of pairwise correlation. There are some higher values between the features that are the result of label encoding (especially **other_installment_plans** and **housing**). Indeed, those features are naturally correlated by construction and, as a matter of fact, their correlation is negative, since their values are all 0s and 1s, and if one of them has a 0 there will be for sure another one having 1 for that data point. This, together with the fact that each original feature has been encoded in a few number of new ones (from 2 to 4) and that in some cases their values were highly imbalanced, contribute to their correlation values.

In addition, it can be interesting to notice that there is a positive correlation (≈ 0.65) between **amount** and **duration** that can be explained by the fact that generally it needs more time to pay an higher credit, and between

property and **housing_3** (≈ 0.52), since the latter is 1 when the person has its own house, and of course this feature's meaning partially overlaps with the other one.

In any case, there are not situations where some features can be considered duplicated. Also by looking at the dendrogram [Figure 12] (built by hierarchically clustering the pairwise features' distances using the average linkage), it can be seen that even the nearest features are quite distant among them.


Figure 11: Pairwise correlation according to Pearson coefficient

5.3   PCA

Principal Component Analysis is one of the most popular and simple techniques of dimensionality reduction. Its goal is to find a linear mapping that project the m data points in a low dimensional space, from d features to n, with $n << d$, in such a way that the information lost is as small as possible. This technique has two formulations that lead to the same solution from different points of view.

- the **minimum error** formulation, that finds a solution to the problem

$$\operatorname*{argmin}_{U \in \mathbb{R}^{n,d}} \sum_{i}^{m} \left\| \vec{\mathbf{x}}_i - UU^T \vec{\mathbf{x}}_i \right\|_2^2, \quad \text{subject to } U^T U = \mathbb{I} \tag{4}$$

The idea is that we want minimise the error, that is the information we loose, in projecting the data in a lower dimensional space. Indeed, the error is computed between each original point $\vec{\mathbf{x}}$ and its reconstruction $\hat{\vec{\mathbf{x}}} = UU^T \vec{\mathbf{x}}$.

Assuming that the data points are centered ($\vec{\mu} = \vec{\mathbf{0}}$), and given the matrix of data points $X \in \mathbb{R}^{m,d}$, m number of points and d number of features, the solution to find the linear mapping (the orthogonal matrix U^T) is finding the n eigenvectors ($\vec{\mathbf{u}}_1, \vec{\mathbf{u}}_2, ..., \vec{\mathbf{u}}_n$) corresponding to the n biggest eigenvalues ($\lambda_1 \geq \lambda_2 \geq ... \geq \lambda_n \geq ... \geq \lambda_d$) of the covariance matrix of X. Those eigenvectors (the *principal components*) are the column of the matrix U, and the sum of the remaining eigenvalues is the error done in projecting the data in low dimensional space.
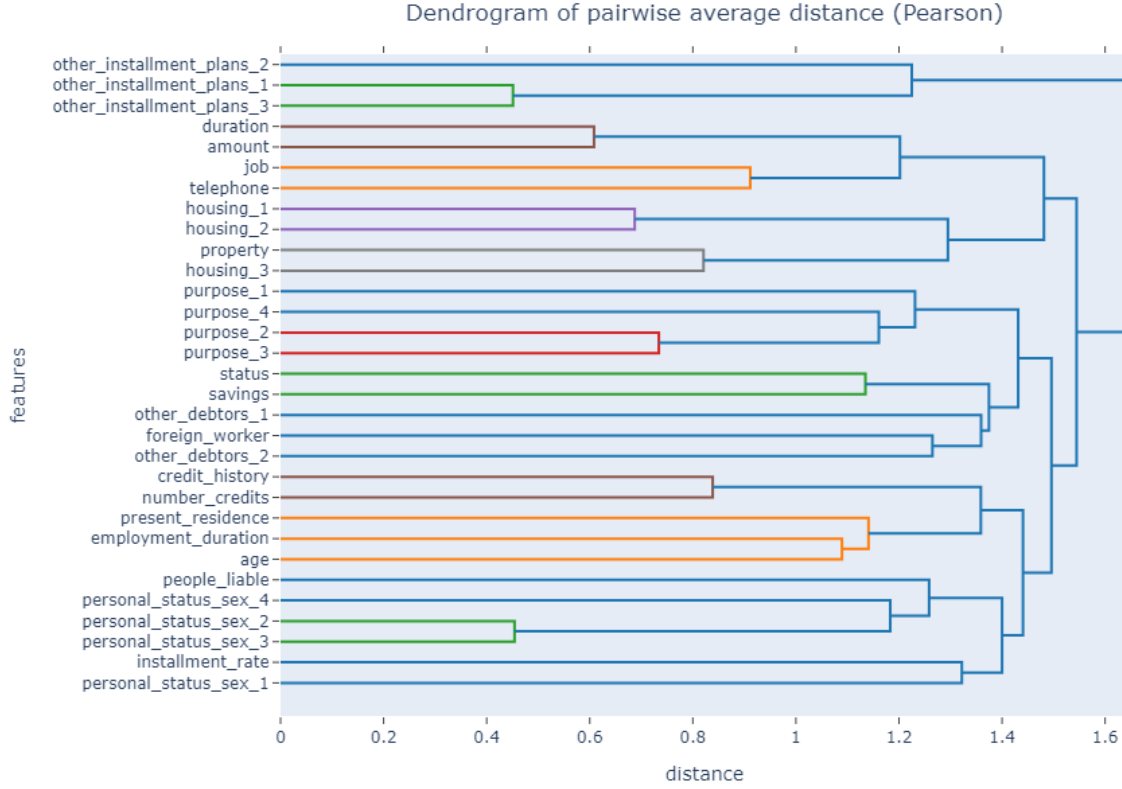
Figure 12: Dendrogram of pairwise distance between features

However, since the data distribution is not known, the covariance matrix is estimated through the scatter matrix, also called sample covariance matrix

$$S = X^T X \tag{5}$$

still assuming $\vec{\mu} = \vec{0}$.

- the **maximum variance** formulation, that aims to find as principal components the directions where the variance of the projection of data points is maximized, since the idea is that the directions where the variance is maximum, are the ones where the information carried also is maximal. Those directions are to be found one at a time, and for each one the variance of data along that direction should be maximized, enforcing the constraints that each principal component must have norm 1 and it has to be orthogonal with respect to the previous ones.

  In the end, it turns out that finding the first n principal components $\vec{u}_1, \vec{u}_2, ..., \vec{u}_n$ means to find the n eigenvectors of the sample covariance matrix (equation 5) and that their corresponding eigenvalues, are exactly the variance of the points projected along that directions. In particular, the first n eigenvectors found with the maximization variance method, are the ones corresponding to the n maximum eigenvalues of the matrix S. Thus, the two methods are exactly equivalent.

In order to decide how many principal components to take, a trade off must be found in order to take neither too little information, nor too many features. Here [Figure 13] both the explained variance (the eigenvalues drown in descending order) and the cumulative explained variance (for each eigenvalue, its ratio w.r.t. the sum of all eigenvalues is computed, then drown in a cumulative way) are showed. In this case, since there is an elbow at 15 principal components, and since at that point the cumulative explained variance is sightly higher than 85%, I decided to take 15 as number of principal components to keep.

5.4 Kernel PCA

A possible variation of the above technique is its Kernel extension. In fact, the Kernel PCA does the same as PCA, but it operates in a high dimensional space. The idea behind this relies on the intuition that many datasets, which are not linearly separable in their space, can be made linearly separable by projecting them into a higher dimensional
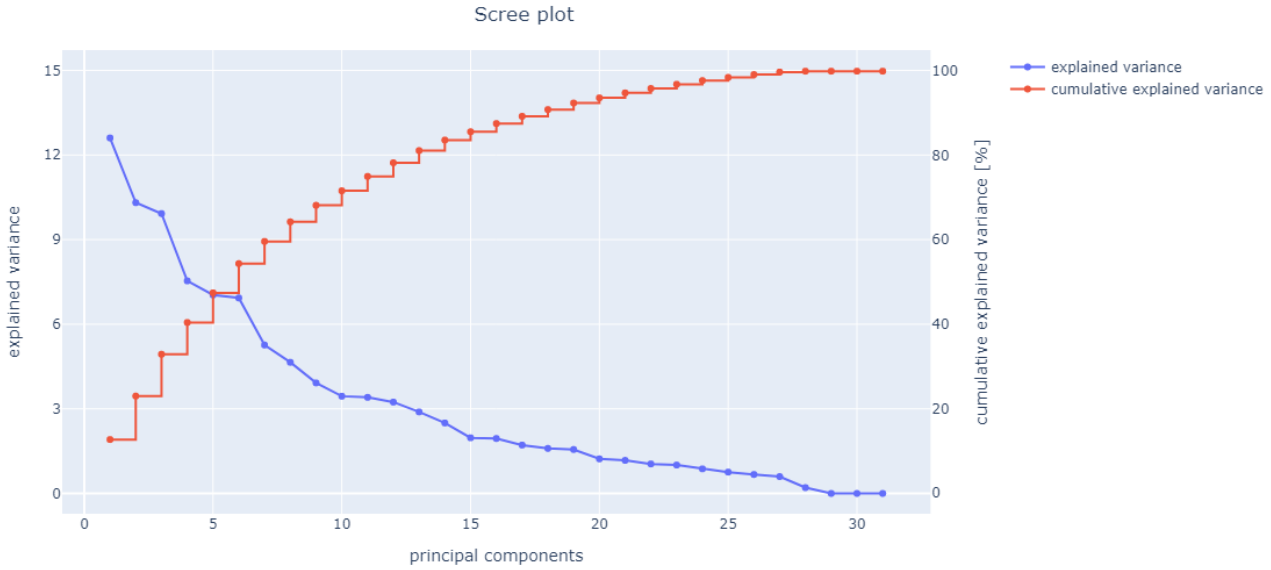
Figure 13: Scree plot of principal components

space; then, in that space, PCA is more likely to give a good result.

However, assuming that a mapping Φ from the feature space to a Reproducing Kernel Hilbert Space (RKHS) exists, there is no need to map all points $\vec{x} \rightarrow \Phi(\vec{x})$ and apply PCA on the new space. Indeed, thanks to the *Mercer's theorem*, we only need a symmetric function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ positive semi-definite to have the assurance that k implements an inner product in some RKHS; thus, there is no need neither to compute Φ, nor to project the data points in the Hilbert space. Specifically, given n data points $(\vec{x}_i)$ and the matrix $K_{n \times n}$ whose element $K(i,j) = \Phi(\vec{x}_i)^T \cdot \Phi(\vec{x}_j) = k(\vec{x}_i, \vec{x}_j)$, the KPCA solution is just finding the top d eigenvectors of K, corresponding to the highest eigenvalues.

In particular, in the implementation by `sklearn`[4], there are four possible non-linear kernels: poly, rbf, sigmoid and cosine. To compare those kernels, a grid search (5-fold cross validation, the default one for `sklearn.model_selection.GridSearchCV`[5] function) has been performed for each one to choose the best hyperparameters that minimize the mean squared error on the reconstruction of data points; in terms of this metric the best one seems to be `poly`, the polynomial kernel. However, by looking at their Scree plots [Figure 14], they seem pretty similar both among them and with respect to the result of standard PCA [Figure 13]: they all have an elbow in explained variance graph around 15 principal components, and in that point the cumulative explained variance is around 85% (82% for `rbf` kernel).

5.5   mRMR

Unfortunately, a downside of [Kernel] PCA, is that the resulting principal components, that will be the new features used in the classification algorithms, are no longer features with a specific meaning. Instead, they are mappings (linear or not) of the original ones, thus making the interpretability of the classification models very difficult. For this reason, it is interesting to explore methods of dimensionality reduction whose goal is to keep a limited number of the original variables with some specific criteria. For example, the algorithm written by Peng *et al.*[3] operates an heuristic feature selection based on minimum Redundancy and Maximum Relevance (mRMR)[6]. Basically it tries to select features by minimizing the dependency among them and maximizing the mean of the mutual information between every feature and the class label. The authors of the paper have investigated two ways to combine redundancy and relevance: the Mutual Information Quotient (MIQ) and the Mutual Information Difference (MID); that is using respectively their ratio or their difference. To evaluate the classification results of this method with the others , it has been set to select 15 features, that for the two schemas are respectively:

- **MIQ**: status, foreign_worker, other_installment_plans_2, purpose_2, credit_history, property, savings, housing_1, personal_status_sex_1, other_installment_plans_1, other_debtors_1, other_debtors_2, age, number_credits, installment_rate,

---

[4]documentation at: `https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.KernelPCA.html`

[5]documentation at: `https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html`
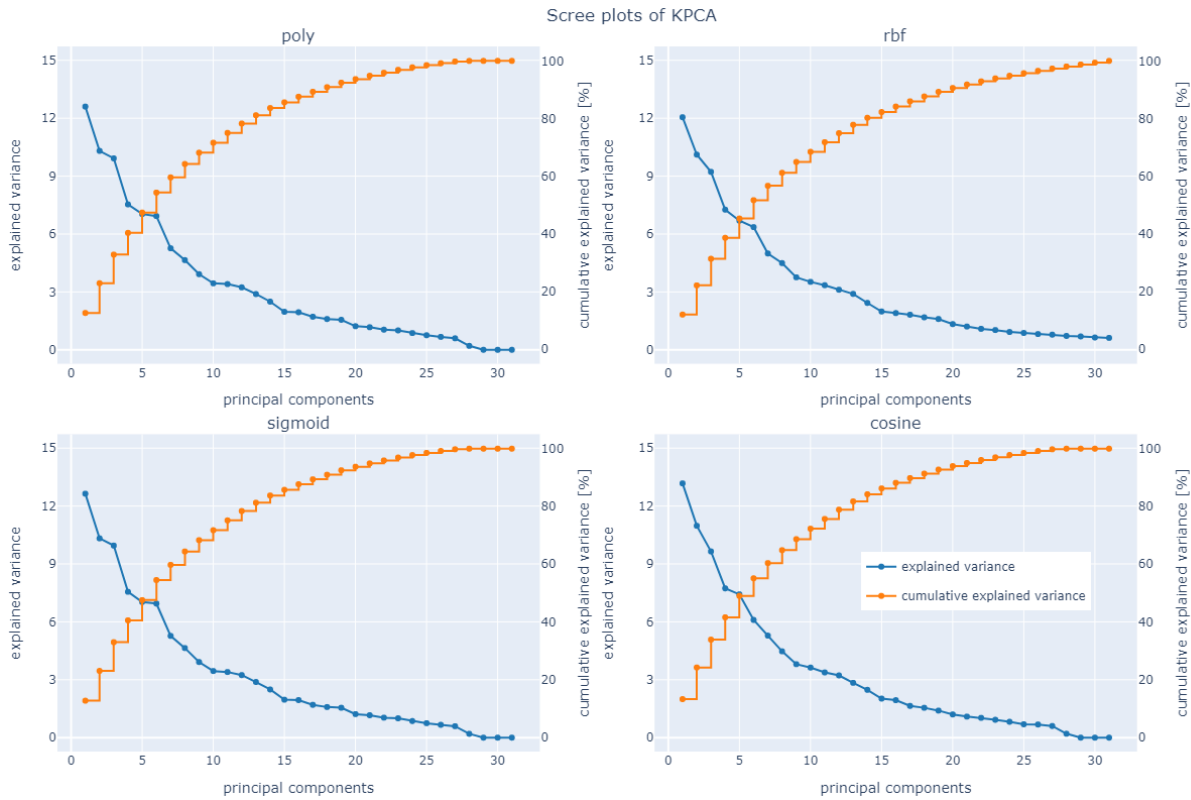
[6]documentation at: `https://pypi.org/project/pymrmr/`

Figure 14: Scree plot of KPCA with four different kernels

- **MID**: status, foreign_worker, credit_history, purpose_2, savings, housing_2, other_installment_plans_3, personal_status_sex_1, other_debtors_1, amount, duration, number_credits, other_debtors_2, age, installment_rate

As can be seen, most of the features selected with the two methods are the same.

6 Cross validation pipeline

At this point, the classification models can be trained to be able to detect good and bad creditors. However, each model can depend on one or more hyperparameters, that are the parameters that control the training process in the machine learning algorithm. In order to choose the best values that make the models to outperform in the desired task, a cross validation process is needed.

First of all, 5 different training sets are prepared, one for each dimensionality reduction technique discussed above (Section 5): PCA, KernelPCA, mRMR with method MIQ and mRMR with method MID, plus the one without dimensionality reduction.

Then, for each of those training sets, a stratified 10-fold cross validation (`GridSearchCV`[7] function of `sklearn` library) is run for each combination of dataset balancing technique (Section 7) and classification algorithm (Section 8). In particular, a few comments are necessary:

- the 10-fold has to be stratified because at each step the validation fold, used to test performances, must have the same characteristics of the test set;

- for the same reason, the dataset balancing technique is applied only to the 9 folds used for training;

- unfortunately, the same was not possible for the dimensionality reduction techniques, since the mRMR method has not been implemented to be part of a pipeline, thus they have all been applied to the entire training set (but not to the test set). Assuming that the 10 fold are sampled in a random way, at each step the 10^{th} one should not have such different characteristics with respect to the the other ones;

---

[7]documentation at: `https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html`

- using a k-fold cross validation instead of the classical train-validation split of the training set, is helpful in trying to reduce overfitting.

During the cross validation step, the chosen metric to be maximized is the F1 score. Given precision $= \frac{\text{TP}}{\text{TP+FP}}$ and recall $= \frac{\text{TP}}{\text{TP+FN}}$:

$$F1 = 2 * \frac{\text{precision} * \text{recall}}{\text{precision} + \text{recall}} \tag{6}$$

It is probably the best score for this kind of problems, because the accuracy may not be the good choice when dealing with an imbalanced dataset, since it just counts the correctly classified points with respect to all the points, without taking into account the label imbalance. The F1 score, instead, enforces the models to correctly classify the minority class points.

7 Dataset balancing

Dealing with a dataset where the class labels' distribution is imbalanced [Figure 1] may lead to poor performing classification algorithms, because most of them assume that class labels are equally distributed and that an be problematic mostly in cases like this one, where we are more interested in detecting the minority class than the majority one: we want the algorithms to be able to properly detect the bad creditors.
For this reason, three main techniques have been investigated: oversampling, undersampling and SMOTE.

7.1 Undersampling

With this technique, given a training set with N points in the minority class and M in the majority one $(M > N)$, in order to balance the class labels' distribution, N points are randomly sampled from the majority class. The main downside is that we are removing samples from the original dataset, possibly taking away useful information to correctly classify data.

7.2 Oversampling

This technique aims to reach a uniform distribution in class labels by randomly sampling data from the minority class, until its number of samples matches the one of the majority class. Of course there is no longer the problem of loosing information, however, given the big gap between the two labels, overfitting is more likely to occur. Especially with this dataset, whose minority class has already been oversampled from 5% to 30% as reported by Grömping [7].

7.3 SMOTE

Synthetic Minority Oversampling Technique (SMOTE) is again an oversampling technique, but instead of randomly sampling points from the minority class, they are linearly combined to synthetically generate new points. In order to augment data, for each data point in the minority class, its k-nearest neighbours (e.g. k= 5) are found and, among them, one is randomly selected. The new point is generated as a convex combination of the two data points.
Although this can be a good way to balance the dataset without incurring in overfitting, it is designed to work only with continuous data. Indeed categorical features, even if properly encoded, have a fixed and limited amount of possible values; thus, it is likely that synthetically generated points would have inconsistent values.

8 Classification models

In this section, all the classification models trained with this dataset are investigated. As mentioned before, a 10-fold cross validation is run to decide the best dimensionality reduction technique, balancing technique as well as the models' hyperparameters.

8.1 Decision Tree

Decision tree is a very popular tool used both for classification and regression. At the basis, the idea is to divide the predictor space into N non overlapping regions where the assumption is that in every region only one label is assigned, such that the prediction error is minimized with the lowest number of regions. Unfortunately, this procedure is computationally unfeasible, because the minimum amount of regions involved in the minimization process is equal to the number of training points. Thus, decision trees are built with a top-down greedy procedure, where at each

Figure 15: Gini index compared to cross-entropy. Image by [9]

step the best current split is done and so on until in each region there is only one label. To decide which one is the best, a criterion must be used, either the *Gini index* (eq 7), a measure of node purity that assumes small values when a node contains more observations of the same class, or the *cross-entropy* (eq 8), very similar to the previous one, but it tends to penalize less small impurities, since it assumes lower values with respect to *Gini* [Figure 15]; here it is treated as an hyperparameter of the model (`criterion`).

Given K the number of classes (here 2), m the current split (or region) and $\hat{p}_{mk}$ the proportion of the points from the k^{th} in the m^{th} region:

$$\text{Gini index} = \sum_{k=1}^{K} \hat{p}_{mk}(1 - \hat{p}_{mk}) \tag{7}$$

$$\text{Cross-entropy} = \sum_{k=1}^{K} \hat{p}_{mk} \log \hat{p}_{mk} \tag{8}$$

However, the main disadvantage of this model is that it is prone to overfitting. Indeed, if the tree is built until no further split can be done, the model ends up describing the training set itself without a predicting goal. For this reason, two main strategies have been applied:

- one way is to set some hyperparameters of the problem to limit this phenomena, such as the depth of the tree (`max_depth`), the minimum number of samples in leaves (`min_samples_leaf`) or the minimum number of samples to split a node (`min_samples_split`). They can be expressed both in absolute values or as a fraction of the whole training set.

- another way is to do the other way around: the tree is firstly built until every split can be done, then it is pruned by progressively substituting subtrees with leaves. The objective is to minimize the error function as before, but it now has a regularization term that consists in the number of leaves of the node weighted for an hyperparameter α (`ccp_alpha`).

8.2   Tree-based ensemble methods

Another way to overcome overfitting and variance and improve generalization of decision tree classifier, is by applying an ensemble method. The idea is that combining multiple weak models (trees in this case) can improve results with respect to keeping a single model. In classification, the final prediction of a new point is the majority of the models' predictions. Here two procedures are taken into account, in both cases, they have been cross-validated and then trained with the strategies discussed above (with and without pruning); however to perform a fair comparison, decision tree hyperparameters have not been tuned again, but the ones previously obtained are used.

8.2.1   Bagging

With this method, B sets of point are generated by taking samples (with repetition) from the original training set (this procedure is called *Bootstrap*), then from each one a tree is built as explained before. The tuned hyperparameters are: the number of samples of each set of point (`max_samples`) expressed as the percentage of the training set, and the number of estimators (`n_estimators`), that is also the number of sets B.

8.2.2 Random Forest

Random forest, instead, is an improvement of bagging. Indeed, it builds decision trees on bootstrapped training samples, but in addition, each time a split has to be done in the tree building procedure, only a limited and random selection of features is considered. Here only the number of trees in the random forest is tuned (`n_estimators`).

8.3 Logistic Regression

Logistic regression is a classification method based on the composition of the sigmoid function σ and the class of linear functions. Thus, the hypothesis classifier is

$$H = \{\vec{x} \mapsto \sigma(\langle \vec{x}, \vec{\beta} \rangle) : \vec{\beta} \in \mathbb{R}^d\} \tag{9}$$

Where $H(\vec{x})$ is a real value in the interval $(0,1)$ and it corresponds to the probability of the point $\vec{x}$ to have label $y = 1$ (viceversa, the probability to have label $y = 0$ is $1 - H$). However, since in classification we want the model to predict a label, a threshold is needed to decide whether the outcome is label 0 or 1.

In this model, the hyperparameters to be tuned involve the regularization term whose role is to avoid overfitting and simplify the model itself. Here two kind of regularizations are considered (`penalty`): the Ridge regression ($\|\vec{\beta}\|_2^2$), or the LASSO one($\|\vec{\beta}\|_1$); both are weighted for $\lambda = \frac{1}{C}$ (in this case `C` is the hyperparameter tuned in cross-validation).

8.4 Support Vector Machine

8.5 K-Nearest Neighbours

-logistic regression (also + regularization ridge or lasso) -tree -random forest (less simple to explain results) -svm (linear, rbf, polinomial, sigmoid) -knn -fda -qda? -lda? bayes? -eventually simple mlp

9 Classification results

-per ogni combinazione (dim-reduction - algoritmo) fare ROC, accuracy, confusion matrix, recall, precision, F1 (armonic mean between precision and recall) -istogrammi ?

References

[1] P. C. Mahalanobis. "On the generalised distance in statistics". In: *Proceedings of the National Institute of Sciences of India 2*. Vol. 1. 1936, pp. 49–55.

[2] H. Hofmann. *Statlog (German Credit Data) Data Set*. 1994. URL: `https://archive.ics.uci.edu/ml/datasets/Statlog+%28German+Credit+Data%29`.

[3] Hanchuan Peng, Fuhui Long, and Chris Ding. "Feature selection based on mutual information: criteria of max-dependency, max-relevance, and min-redundancy". In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 27 (2005), pp. 1226–1238. URL: `http://home.penglab.com/papersall/docpdf/2005_TPAMI_FeaSel.pdf`.

[4] Shai Shalev-Shwartz and Shai Ben-David. *Understanding Machine Learning: From Theory to Algorithms*. USA: Cambridge University Press, 2014. ISBN: 1107057132.

[5] Dheeru Dua and Casey Graff. *UCI Machine Learning Repository*. 2017. URL: `http://archive.ics.uci.edu/ml`.

[6] Mehryar Mohri, Afshin Rostamizadeh, and Ameet Talwalkar. *Foundations of Machine Learning*. 2nd ed. Adaptive Computation and Machine Learning. Cambridge, MA: MIT Press, 2018. 504 pp. ISBN: 978-0-262-03940-6.

[7] U. Grömping. *South German Credit Data: Correcting a Widely Used Data Set*. Tech. rep. 4/2019, Reports in Mathematics, Physics and Chemistry, Berlin: Department II, Beuth University of Applied Sciences, Apr. 2019. URL: `http://www1.beuth-hochschule.de/FB_II/reports/Report-2019-004.pdf`.

[8] U. Grömping. *South German Credit (UPDATE) Data Set*. 2020. URL: `http://archive.ics.uci.edu/ml/datasets/South+German+Credit+%28UPDATE%29`.

[9] Stathis Kamperis. *Decision Trees: Gini index vs entropy*. Apr. 2021. URL: `https://ekamperi.github.io/machine%20learning/2021/04/13/gini-index-vs-entropy-decision-trees.html`.

[10] Sergen Cansiz. *Multivariate Outlier Detection in Python*. URL: `https://towardsdatascience.com/multivariate-outlier-detection-in-python-e946cfc843b3`.

[11] numpy documentation. *numpy.cov*. URL: `https://numpy.org/doc/stable/reference/generated/numpy.cov.html`.

[12] Plotly. *Box Plots in Python*. Accessed: 2021-08-09. URL: `https://plotly.com/python/box-plots/`.

[13] Wikipedia, the free encyclopedia. *Box plot*. Last edit: 2021-07-07 Accessed: 2021-08-09. URL: `https://en.wikipedia.org/wiki/Box_plot`.

[14] Wikipedia, the free encyclopedia. *Deutsche Mark*. Last edit: 2021-07-10 Accessed: 2021-08-09. URL: `https://en.wikipedia.org/wiki/Deutsche_Mark`.